



Post hoc train-neighbor similarity audit for fixed sequence benchmarks

1 Fixed benchmark



single score; neighbor context hidden

2 Similarity audit



Max train-neighbor
k-mer Jaccard
within length window

Thresholds

30% 40% 60% 80% 100%



no retraining


fixed
predictions

3 Similarity-aware report

Benchmark
score

mAP = 0.888

Profile



span = 11.7 pp

benchmark score + profile

Separates low-similarity generalization
from family expansion